


```
import pandas as pd

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
import xgboost as xgb
from sklearn.metrics import accuracy_score, classification_report

# Load the dataset, assuming it's a CSV despite the .xls extension
try:
    df = pd.read_csv('/content/WA_Fn-UseC_-Telco-Customer-Churn.xls')
    print("File loaded successfully as CSV.")
except Exception as e:
    print(f"Error loading file as CSV: {e}")
    print("Attempting to load as Excel with openpyxl engine (for modern .xlsx")
    try:
        df = pd.read_excel('/content/WA_Fn-UseC_-Telco-Customer-Churn.xls', er
        print("File loaded successfully as Excel.")
    except Exception as e_excel:
        print(f"Error loading file as Excel: {e_excel}")
        print("Could not load the file. Please ensure it's a valid CSV or Excel")
        exit() # Exit if file cannot be loaded

# Display the first few rows and info
print(df.head())
print(df.info())

# --- Data Preprocessing for XGBoost ---

# Drop 'customerID' column as it's not useful for prediction
df = df.drop('customerID', axis=1)

# Convert 'TotalCharges' to numeric. It might have spaces or non-numeric values
df['TotalCharges'] = pd.to_numeric(df['TotalCharges'], errors='coerce')
# Handle missing 'TotalCharges' values (if any) by filling with the mean or median
df['TotalCharges'] = df['TotalCharges'].fillna(df['TotalCharges'].mean())

# Convert 'Churn' target variable to numerical (Yes=1, No=0)
df['Churn'] = df['Churn'].apply(lambda x: 1 if x == 'Yes' else 0)

# Identify categorical and numerical features
categorical_features = df.select_dtypes(include='object').columns
numerical_features = df.select_dtypes(include=['int64', 'float64']).columns.dri

# Apply One-Hot Encoding to categorical features
df_encoded = pd.get_dummies(df, columns=categorical_features, drop_first=True)

# Separate features (X) and target (y)
X = df_encoded.drop('Churn', axis=1)
```

```

X = df_encoded.groupby('Churn', axis=1)
y = df_encoded['Churn']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# --- XGBoost Model Training ---

# Initialize XGBoost Classifier
xgb_model = xgb.XGBClassifier(
    objective='binary:logistic', # For binary classification
    eval_metric='logloss',       # Evaluation metric
    n_estimators=100,           # Number of boosting rounds
    learning_rate=0.1,          # Step size shrinkage
    max_depth=3,                # Maximum depth of a tree
    random_state=42
)

# Train the model
xgb_model.fit(X_train, y_train)

# --- Model Evaluation ---

# Make predictions on the test set
y_pred = xgb_model.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
report = classification_report(y_test, y_pred)

print(f"\nAccuracy: {accuracy:.4f}")
print("\nClassification Report:")
print(report)

```

File loaded successfully as CSV.

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	
0	7590-VHVEG	Female		0	Yes	No	1	No
1	5575-GNVDE	Male		0	No	No	34	Yes
2	3668-QPYBK	Male		0	No	No	2	Yes
3	7795-CFOCW	Male		0	No	No	45	No
4	9237-HQITU	Female		0	No	No	2	Yes

	MultipleLines	InternetService	OnlineSecurity	...	DeviceProtection	\
0	No phone service		DSL	No	...	No
1		No	DSL	Yes	...	Yes
2		No	DSL	Yes	...	No
3	No phone service		DSL	Yes	...	Yes
4		No	Fiber optic	No	...	No

	TechSupport	StreamingTV	StreamingMovies	Contract	PaperlessBilling	\
0	No	No	No	Month-to-month	Yes	
1	No	No	No	One year	No	

#	tenure	PhoneService	Contract	Month-to-month	One year	Yes
2	No	No	No	Month-to-month		Yes
3	Yes	No	No		One year	No
4	No	No	No	Month-to-month		Yes

#	PaymentMethod	MonthlyCharges	TotalCharges	Churn
0	Electronic check	29.85	29.85	No
1	Mailed check	56.95	1889.5	No
2	Mailed check	53.85	108.15	Yes
3	Bank transfer (automatic)	42.30	1840.75	No
4	Electronic check	70.70	151.65	Yes

[5 rows x 21 columns]

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 7043 entries, 0 to 7042

Data columns (total 21 columns):

#	Column	Non-Null Count	Dtype
0	customerID	7043	non-null
1	gender	7043	non-null
2	SeniorCitizen	7043	non-null
3	Partner	7043	non-null
4	Dependents	7043	non-null
5	tenure	7043	non-null
6	PhoneService	7043	non-null
7	MultipleLines	7043	non-null
8	InternetService	7043	non-null
9	OnlineSecurity	7043	non-null
10	OnlineBackup	7043	non-null
11	DeviceProtection	7043	non-null
12	TechSupport	7043	non-null
13	StreamingTV	7043	non-null
14	StreamingMovies	7043	non-null
15	Contract	7043	non-null
16	PaperlessBilling	7043	non-null
17	PaymentMethod	7043	non-null
18	MonthlyCharges	7043	non-null
19	TotalCharges	7043	non-null
20	Churn	7043	non-null

dtypes: float64(1), int64(2), object(18)

memory usage: 1.1+ MB

